



OpusPocus

NMT Training Pipeline Manager

August 27

varis@ufal.mff.cuni.cz

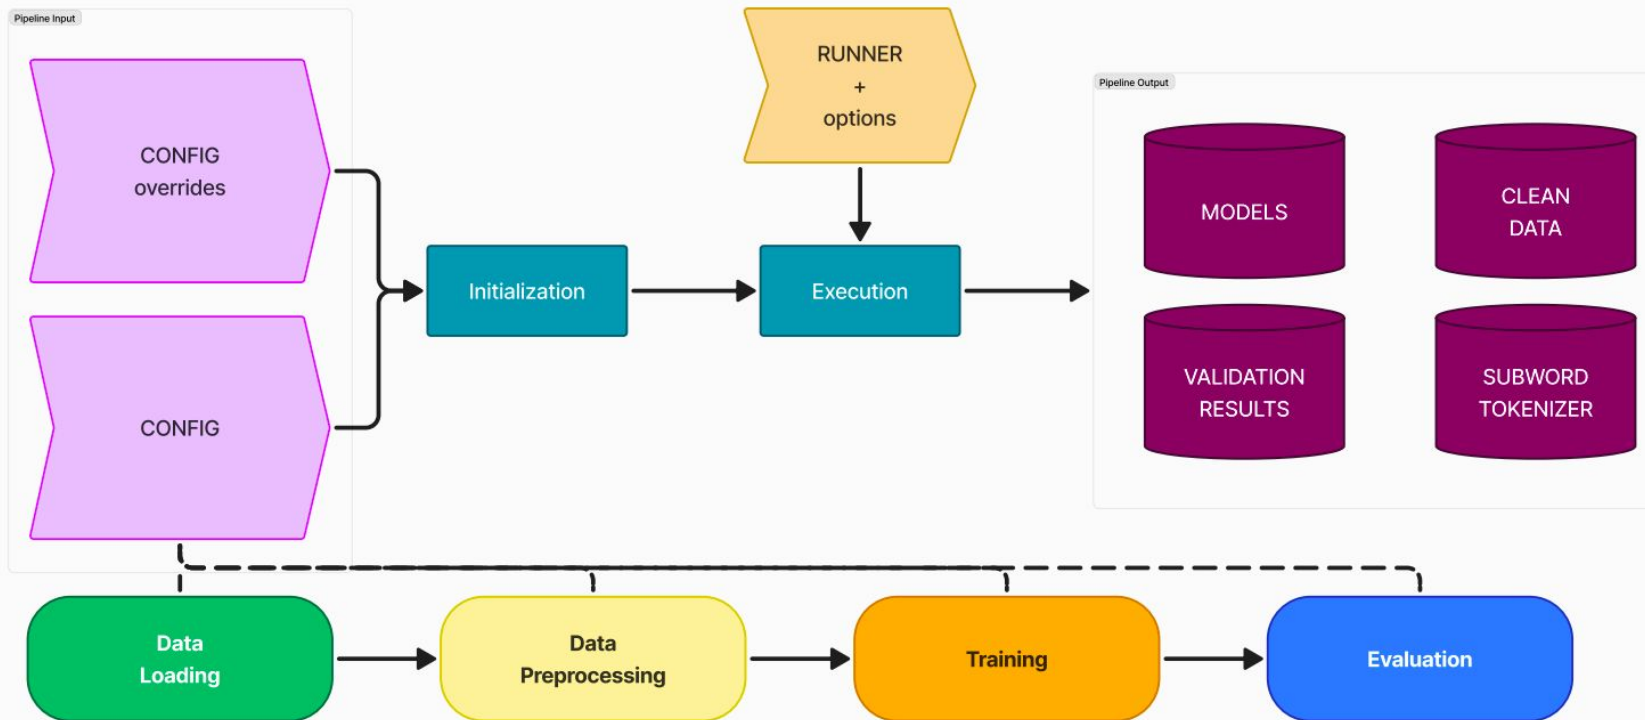
Quick Introduction

Motivation



- NMT Training - too many stages to execute manually -> automation
- Existing solutions
 - Makefile, SnakeMake
 - too general,
 - hard to debug/maintain (when combined with HPC schedulers),
 - known issues on HPC (LUMI)
 - handling resubmission after time-limit
 - effective pipeline recovery after sharded task failure
- OpusPocus
 - focused only on NMT
 - in future could be extended to other NLP tasks
 - developed with MT community needs in mind
 - robust with regards to HPC (Slurm)

Basic Workflow





- Experiment replicability:
 - config describes full train pipeline (from data acquisition up to evaluation)
- Adaptation to many language (language pairs)
 - executing overwritten pipeline config
- Experiments with data preprocessing
 - Variations of the preprocessing pipeline + fixed model training pipeline
- Testing model variants
 - backtranslation
 - data mixing strategies
 - iterative backtranslation

Exercises

Installation



Python \geq 3.9 (tested with Python 3.10.7)

(Use virtualenv or Conda to create a clean Python installation)

1. Clone the git repository (mtm2025 branch) and install dependencies:

```
$ git clone -b mtm2025 git@github.com:hplt-project/OpusPocus.git
```

```
$ cd OpusPocus
```

```
$ pip install -r requirements.txt
```

2. Download data:

```
$ mkdir data
```

```
$ scripts/get_data_hplt_v2.0.py --languages eu ca --data-dir data/
```

Installation (cont.)



Install Marian (ideally GPU version):

```
$ git clone git@github.com:marian-nmt/marian.git marian
$ mkdir marian/build && cd marian/build
$ cmake .. \
    -DCMAKE_BUILD_TYPE=Release \
    -DCOMPILER_CUDA=ON \
    -DCUDA_TOOLKIT_ROOT_DIR=<CUDA_ROOT_LOCATION> \
    -DUSE_CUDNN=ON \
    -DCUDNN_LIBRARY=<libcudnn.so_LOCATION> \
    -DCUDNN_INCLUDE_DIR=<CUDNN_INCLUDE_LOCATION> \
$ make -j 8
```

(At the end, the Marian should be located at OpusPocus/marian)

Basic Execution



\$./go.py <subcommand> [options]

- Subcommands:
 - **run** - executes pipeline
 - **stop** - stop a running pipeline
 - **status** - print pipeline status
 - **traceback** - step dependency structure with their status

\$./go.py <subcommand> **--help**

- lists subcommand options

Example 0: Running Minimal Pipeline



```
$ ./go.py run --pipeline-config config/00-pipeline.minimal.yml
```

1. No data processing
2. Single direction model training + evaluation

Check status:

```
$ ./go.py status --pipeline-dir experiments/en-eu/pipeline.minimal
```

Stop the pipeline:

```
$ ./go.py stop --pipeline-dir experiments/en-eu/pipeline.minimal
```

Continue execution:

```
$ ./go.py run --pipeline-dir experiments/en-eu/pipeline.minimal  
--reinit-failed
```

Pipeline Config



- YAML format, processed with OmegaConf
 - easy to read, simple processing in Python
 - variable interpolation support
 - `${full.path.to.variable}` (e.g. `${pipeline.pipeline_dir}`)
 - or relative paths
 - `${.current_level_var}`
 - `${..level_above_var}`
 - `${...two_levels_above_var}`

Supports CLI overwriting:

```
$ ./go.py run --pipeline-config config/00-pipeline-minimal.yml  
pipeline.src_lang=$lang pipeline.tgt_lang=en
```

Pipeline Config Structure



- pipeline:
 - **.steps** - list of steps and their dependencies
 - **.targets** - final pipeline steps to be executed
 - also implies execution of the target step dependencies
 - other variables - used as defaults for non-specified step-arguments
- runner:
 - **.runner** - name of the runner
 - **.runner_resources** - global resources for runner execution
 - other runner specific options
- step (pipeline.steps[i]):
 - list of step argument definitions
 - each step having a separate *dict* of arguments

Step Config Structure



- OpusPocusStep (general args):
 - **step** -
 - **step_label** -
 - **runner_resources** (dict) -
 - ***_step** - step-specific dependencies
 - represented by step_label of a given dependency
- CorpusStep derivatives:
 - take previous (corpus) step and apply a transformation to that corpus
 - args:
 - **src_lang** - (source) corpus language (required)
 - **tgt_lang** - target corpus language (optional)
 - **prev_corpus_step** - step containing corpus being transformed
- Default values:
 - taken from pipeline.<argument> name if available

Pipeline Directory



Created during **INITIALIZATION** (before execution)

- **pipeline_root/**
 - **pipeline.config** - configuration of the pipeline
 - **<step_label>/** - subdirectory containing the <step_label> step details
- **<step_label>/:**
 - **output/** - step-produced output
 - **logs/** - log files
 - **temp/** - temporary (input) files - deleted when step is **DONE**
 - **step.state** - current step execution state
 - **step.parameters** - step initialization parameters
 - **step.dependencies** - list of step dependencies
 - **step.command** - step executable
 - **runner.step_info** - information about a runner submission (job_id, ...)

Pipeline/Step States



Step States

- null (not created yet)
- **INITED** - initialized, step directory created
- **SUBMITTED** - submitted for execution by a RUNNER
- **RUNNING** - currently being executed
- **DONE** - execution 👍
- **FAILED** - execution 👎

PipelineStates similar + INIT_INCOMPLETE and PARTIALLY_DONE

Example 1: Bidirectional Translation Pipeline



```
$ ./go.py run --pipeline-config config/01-pipeline.bidirectional.yml
```

1. Add the opposite translation directions
2. Simplify config - move selected step arguments to global pipeline arguments

Global Pipeline variables



```
pipeline:
  src_lang: en
  tgt_lang: eu
  steps:
    - step: raw
      step_label: test1
    - step: raw
      step_label: test2
```

Equals to:

```
pipeline:
  steps:
    - step: raw
      step_label: test1
      src_lang: en
      tgt_lang: eu
    - step: raw
      step_label: test2
      src_lang: en
      tgt_lang: eu
```

Example 2: Execution on Slurm



```
$ ./go.py run --pipeline-config config/02-pipeline.bidirectional.slurm.yml
```

1. Change the default runner to “slurm” runner
2. Specify the available *runner_resources* for the runner

Runner Resources



- Runner-agnostic representation of the available execution resources
 - # gpus, # cpus, memory size, etc.
- Each runner implements conversion to its own representation
 - “runner_resources.gpus=4” == (Slurm) => “--gpus=4” option
- Each resource is also passed as OpusPocus environment variable to the executed script
 - “runner_resources.gpus” == \$OPUSPOCUS_gpus

Example 3: Step-specific Execution Resources



```
$ ./go.py run --pipeline-config config/03-pipeline.bidirectional.resources.yml
```

1. We want to use different resources for different pipeline steps
 - step-specific *runner_resources* overwrite “global” *runner.runner_resources*

Example 4: Data Preprocessing



```
$ ./go.py run --pipeline-config config/04-pipeline.data_preprocess.yml
```

- preprocess input training data
 - a. clean data with OpusCleaner
 - b. decontaminate (remove training examples similar to valid/test data)
 - c. “gather” corpora into a single corpus file

OpusCleaner – Data Filters



- must be located in the same directory as the corpus files
 - e.g: *data/my_corpus.en-eu.en.gz*
=> e.g.: *data/my_corpus.en-eu.filter.json*
 - no *.filter.json* file => cleaning is skipped
- example filters:
 - *config/opuscleaner_filters/*
 - copy the filters to your raw data directory (*data/en-eu/raw/v2*)

Can be created/edited:

- by directly editing the *.filter.json* files
- via OpusCleaner web UI
 - see <https://github.com/hplt-project/OpusCleaner>

Example 5: NMT with Backtranslation



```
$ ./go.py run --pipeline-config config/05-pipeline.backtranslation.yml
```

- train a mock model with BT data using the preprocessed data from the previous example
 - the previous pipeline must be in the **DONE** state
- “fake” monolingual data (using the parallel corpora instead)
- *auth* and *synth* data mixing via OpusTrainer

OpusTrainer – Configuration



- default config created by OpusPocus
 - limited expression
- more robust configuration directly via OpusTrainer config file:
 - see <https://github.com/hplt-project/OpusTrainer>
 - dataset scheduling, interleaving, curriculum learning, etc.
 - noise introducing schemes for better robustness
 - not dependent on MarianNMT
 - but tested mostly with MarianNMT

Future Work



- Multi-lingual training support
- Support for other training/decoding frameworks
 - HuggingFace
 - NLLB
 - LLM training (?)
- Integrating OpusDistillery for student model training
 - see <https://github.com/Helsinki-NLP/OpusDistillery>
- Fixing issues, improving user experience
 - check <https://github.com/hplt-project/OpusPocus/issues>